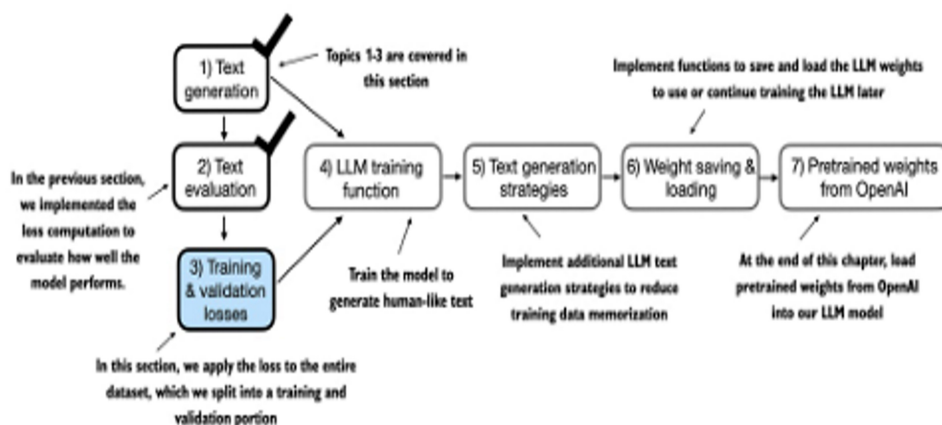


GPT2 Model : Calculating the training and validation set losses



The Verdict

Edith Wharton



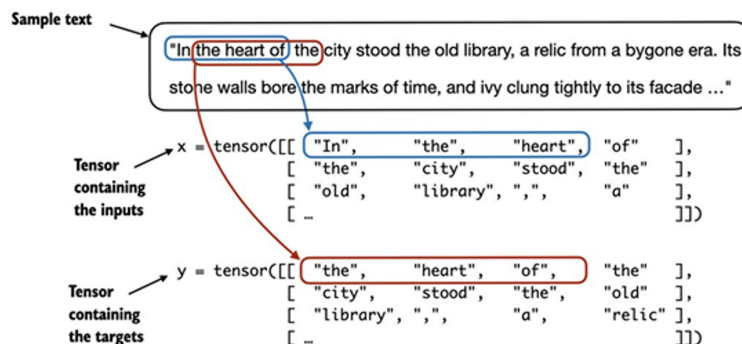
1908

Tokenize (approx 5000 tokens)

Divide the dataset into training and Validation set


```
train_ratio = 0.90
split_idx = int(train_ratio * len(text_data))
train_data = text_data[:split_idx]
val_data = text_data[split_idx:]
```

Use dataloader to chunk training and validation data into input and output datasets.

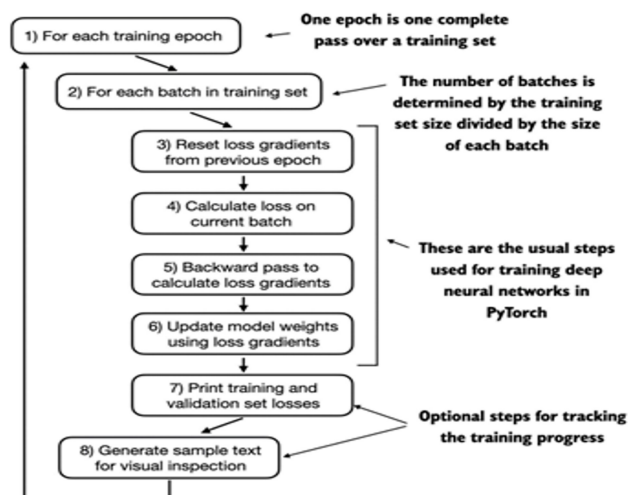


Validation

Training



Pretraining of LLM Model



1) Embedding Parameters : $50257 \times 768 + 1024 \times 768 = 38.4 \text{ M}$

2) Multi Head Attention : $Q, K, V : 3 \times 768 \times 768 = 1.77 \text{ M}$

Output Head = $768 \times 768 = 0.59 \text{ M}$

3) Feed Forward Neural Network : $768 \times 4 \times 768 + 768 \times 4 \times 768 = 4,718,592$

4) Total 12 Transformer blocks : $12 \times (1.77 + 0.59 + 4.72) = 84.2 \text{ M}$

5) Softmax = $768 \times 50257 = 38597376$

Hence Total Parameters = approx 161 M

Why is it 124M and not 164M?

The model's parameter count might be different from the idealized calculation for a couple of reasons:

1. **Tying of the Embedding and Output Layer:** In some implementations of GPT, the input and output embedding layers are tied, meaning that the parameters for the output embedding are not counted separately. This would reduce the total number of parameters.

- If the output embedding is tied to the input embedding, you effectively reduce the embedding parameter count by $V \times d_{\text{model}}$, or $50,000 \times 768$.

- Reducing this from the previous embedding parameters:

$$76,800,000 - (50,000 \times 768) = 76,800,000 - 38,400,000 = 38,400,000$$
$$38,400,000 + 76,800,000 - (50,000 \times 768) = 76,800,000 - 38,400,000 = 38,400,000$$

The total parameter count with tied embeddings:

$$\text{Total Parameters with tied embeddings} = 38,400,000 + 28,311,552 + 56,623,104 = 123,334,656$$
$$\text{Total Parameters with tied embeddings} = 38,400,000 + 28,311,552 + 56,623,104 = 123,334,656$$

2. **Optimizations and Regularizations:** There could be further optimizations or regularization techniques in the implementation that reduce the number of parameters slightly compared to the full theoretical number.

Thus, the model ends up with about **124M** parameters when rounding to the nearest million, instead of the calculated **164M**. The exact difference is mainly due to specific architectural choices such as tying the embeddings and other model design decisions.